1. To what file do you add the tensorflow lite dependency when building an Android app? **1 / 1 point**

- ○ aar.gradle
- ◉ build.gradle
- ○ gradle.build
- ○ build.aar

✓ **Correct**

2. If the Android Neural networks API is available and you want to use it, how would you do that? **1 / 1 point**

- ○ Do nothing, it will work automatically
- ○ Invoke the NNAPI object, and pass the tflite interpreter to it
- ◉ Call the setUseNNAPI method on the interpreter and set its parameter to true
- ○ You can't use the neural networks API with a TensorFlow Lite model

✓ **Correct**

3. If you want to configure the number of threads the interpreter uses, how would you do that? **1 / 1 point**

- ○ Do nothing, it's always single threaded
- ○ Do nothing, it automatically picks the appropriate number of threads
- ○ Call the useThreads() method, and it will apportion the correct number of threads
- ◉ Call setNumThreads and pass it the number of threads you want to use

4. Where's the best place in an Android app to keep your model? **1 / 1 point**

- In the resources folder
- It can really be anywhere, but for consistency use the assets folder
- You don't keep your model in your android App, it should download it at runtime
- In the same folder as the activity that calls it

✓ **Correct**

5. If you tested your converted model and know its valid, but the interpreter cannot load it at runtime on Android, what's the most likely reason? **1 / 1 point**

- You haven't quantized your model
- You have't converted the model to Java or Kotlin format
- You converted your model to iOS format by accident
- You didn't specify that the model should not be compressed in the build.gradle file

✓ **Correct**

6. What is the method signature of the interpreter when you want to do inference? **0 / 1 point**

- interpreter.predict(inputs, predictions)
- predicitons = interpreter.predict(inputs)
- predictions = interpreter.run(inputs)
- interpreter.run(inputs, predictions)

7. What Android data structure is most commonly used to feed image input to the interpreter? 1 / 1 point

- (●) A ByteBuffer
- ( ) A TensorArray
- ( ) A Tensor
- ( ) An Array

✓ **Correct**

8. How many classes of object can a model trained on the COCO dataset recognize? 1 / 1 point

- ( ) 10
- ( ) 800
- ( ) 1000
- (●) 80

✓ **Correct**

9. When performing object recognition, how many dimensions of output tensors are there? 1 / 1 point

- ( ) 10
- ( ) 80
- (●) 4
- ( ) 1

10. How do you get the coordinates of the bounding boxes from the object detection model? 1 / 1 point

- The coordinates are in the first four tensors, read them and simply plot
- The coordinates are in the first tensor, read them and simply plot
- The coordinates are in the first tensor, but arranged differently, you have to sort them before you can plot them
- The coordinates are in tensors 0, 1, 2 and 3

✓ **Correct**